

High Performance Programming

Lecture 1:

”Introduction to Parallel Processing”

Lecturer: Ramoni Adeogun



AALBORG
UNIVERSITET

Content

1. Course introduction
2. Introduction to parallel processing
3. Parallel computing models
4. Effectiveness of parallel processing
5. Summary and Exercises

1. Course introduction

Course information

- **Course Responsible:**

- Ramoni Adeogun, WCN Section (Email: ra@es.aau.dk, Room: FrB 7A3-218)

- Teacher: Stefan Nordborg Eriksen, WCN Section (Email: sne@es.aau.dk, Room: FrB 7A3-215)

- Literatures:

1. Ananth Grama, et. al, Introduction to Parallel Computing, Second edition.
2. Czarnul, P. Parallel Programming for Modern High Performance Computing Systems: Programming with OpenMP, MPI, CUDA, and OpenCL. CRC Press

Others:

1. Xavier, C. (2002). Introduction to Parallel Algorithms. Pearson Education
2. Behrooz P. Introduction to Parallel Processing: Algorithms and Architectures, Kluwer Academic Publishers

Objective and learning outcomes


Objective: To equip students with the skills to design, optimize, and verify high-performance software solutions using parallelism, vectorization, and modern computing architectures like GPUs.

Learning outcomes (from the study regulation)

- **Knowledge**
 - **data structures** used to improve performance
 - basic understanding of **limitations and bottlenecks** in data science solutions
 - **parallelism** and the following issues they raise
 - **vectorization** of operations
 - **GPU-based** operations
 - types of **tests** and their use
 - **quality measures** for the **correctness** of computer science solutions, including: **testing** and **verification**
- **Skills**
 - can resonate and argue for **bottlenecks** in software programs and applications
 - can utilize **parallelism** in the chosen programming language and document the correctness of a given implementation
 - can use and **perform tests in the development process** of a program so that it is documented that its functionality is correct in a number of given cases
 - can use and perform **verification** of simple programs
 - can use correct **professional terminology**
- **Competencies**
 - can solve problems that require high performance by using parallelism in a computer program
 - can argue for the correctness of chosen solutions using tests and verification

Course content and teaching plan (tentative)

- A total of 10 lectures
 - Lecture 1: Introduction to parallel processing
 - Lecture 2: Basic data structures and communication
 - Lecture 3: Paradigms for parallel programs
 - Lecture 4: Analytical modelling of parallel algorithms
 - Lecture 5: Practical aspects of parallel programming
 - Lecture 6: Shared memory parallel programming
 - Lecture 7: Distributed memory parallel programming
 - Lecture 8: GPU programming
 - Lecture 9: Testing and verification
 - Lecture 10: Workshop



Credit: ChatGPT

Examination

- Examination will be **oral**.
- The focus of the examination is to assess you based on the **expected learning outcomes** of the course.
- The exam will cover both the fundamental theoretical concepts and practical aspects of high performance programming.
- Each student is awarded a grade on the **7-point scale**.
- Exact form of the oral exam will be communicated later.


Recommendations

- Prepare for lectures by reading recommended course materials.
- Attend and participate actively during the lectures.
- Solve as many of the exercise problems to enhance your understanding of the concerned concepts.
- Remember to read and understand each problem before attempting to solve it.
- **Remember:** a 5 ECTS course requires approx. **147 hours study time**.
 - Lectures and exercises constitute only about 50 hours.
 - It is your responsibility to appropriately distribute the remainder for reading course materials, preparing for lectures, and examination.

2. Introduction to parallel processing


What is Parallel Processing?

- Parallel processing refers to the **simultaneous** execution of multiple tasks or operations to increase computational efficiency.
 - dividing a problem into smaller sub-problems that can be solved simultaneously.
- multiple processors or cores to work on parts of a problem simultaneously.
- task decomposition and task scheduling to exploit the inherent parallelism of the problem.



Why is parallel processing important?

- Limitations of sequential computing
 - **Moore's Law** slowdown: increase in performance per core has slowed down
 - a single-core processor no longer provides the necessary performance improvements.
 - Sequential algorithms may not scale for large datasets.
- Achieving needed **performance improvement** for large computational problems in applications including:
 - Climate modeling and weather prediction.
 - Computational biology and genomics.
 - Engineering simulations (e.g., fluid dynamics, structural analysis).
 - Real-time data analytics and machine learning.



Parallel processing is the **primary way** to achieve significantly higher application performance for the foreseeable future

Why High-Performance Programming?

1

Higher speed (solve problems faster)

Important when there are “hard” or “soft” deadlines; e.g., 24-hour weather forecast

2

Higher throughput (solve more problems)

Important when we have many similar tasks to perform;
e.g., transaction processing

3

Higher computational power (solve larger problems)

e.g., weather forecast for a week rather than 24 hours,
or with a finer mesh for greater accuracy

From Serial Algorithmic to Parallel Thinking

- So far most or all your courses has assumed that only one thing happens at a time in a program
 - sequential programming → each statement executes in sequence.
- Removing this assumption creates challenges:
 - **Programming:** How can we divide work among threads of execution and coordinate (synchronize) among them?
 - **Algorithms:** How can activities in parallel speed-up a program? → more throughput: work done per unit time
 - **Data structures:** May need to support concurrent access → multiple threads operating on data at the same time.


Short Discussion



Identify aspects of your group project(s) that may benefit from parallel processing?



What do you consider the major show-stoppers for parallelization?



Which of these aspects of parallel processing do you consider important for a computer engineer?

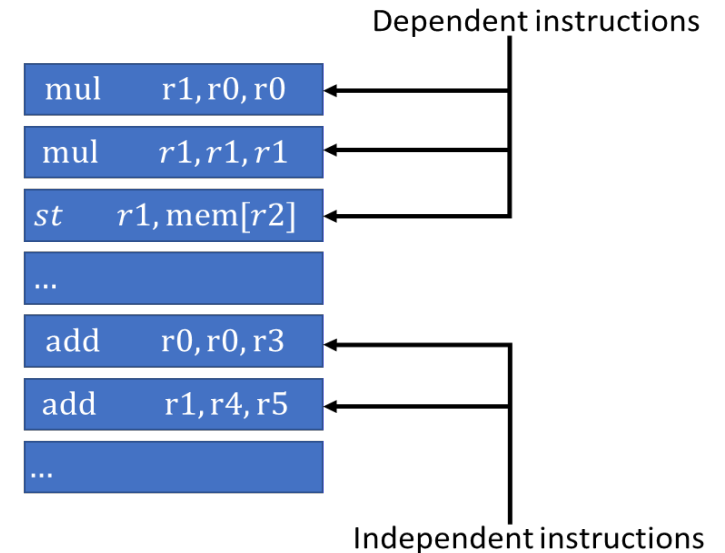
Data structure
Algorithms
Programming

Levels of parallelism

- Levels of parallelism
 - the hierarchical nature or granularity at which computations can execute in parallel.
 - often correspond to hardware and software abstractions.
- 1. Bit-Level Parallelism [1970 to ~1985]
 - Exploits the parallelism inherent in the hardware to process multiple bits of data simultaneously.
 - a 64-bit processor can process more data per clock cycle than a 32-bit processor.
- 2. Instruction-Level Parallelism [~1985 til date]
 - overlaps the execution of multiple instructions within a single processor using techniques like pipelining, out-of-order execution, and superscalar execution.
 - ILP is intrinsic to modern CPU designs and helps improve performance without explicit parallel programming.
- 3. Statement-Level Parallelism
 - Focuses on executing multiple statements of a program concurrently.
 - Compiler optimizations like loop unrolling and vectorization target this level by enabling the parallel execution of statements.
- 4. Task-Level Parallelism
 - Involves splitting a program into distinct tasks that can run independently or concurrently.
 - explicitly controlled by the programmer using threads, processes, or distributed computing paradigms like MPI.



Parallelism in single-core architectures

- What does a processor do? It runs programs
 - Execute instruction referenced by the program counter (PC) → instruction execution will modify machine state: contents of registers, memory, CPU state, etc
 - Move to the next instruction → execute it....
- Processors did in fact leverage parallel execution to make programs run faster, it was just invisible to the programmer
- Instruction level parallelism (ILP)
 - Instructions must appear to be executed in program order.
 - **independent** instructions can be executed
 - simultaneously by a processor without impacting program correctness
 - **Superscalar execution:** processor dynamically finds
 - independent instructions in an instruction sequence
 - and executes them in parallel



Instruction level parallelism (ILP): Example and limitations

- Instruction: $y = \sum_{k=1}^4 x_k \times x_k$



Most available ILP is exploited by a processor capable of issuing four instructions per clock → little gain from using one with more capability

Types of parallelism

- **Task parallelism**


- Different tasks or operations run in parallel, but each task is independent.
- Different components of a program running simultaneously (e.g., data input, calculation, and output).

- **Data parallelism**

- The same operation is applied simultaneously to multiple pieces of data.
- Ex: Matrix operations where each element can be processed in parallel.




- **Pipeline parallelism**

- Tasks are divided into stages, and different stages can be executed concurrently.
- Ex: image processing pipeline → one part of the image is processed by one unit and the next part by another



Parallel programming models



- Shared Memory Model
 - Processors share access to a common memory space, allowing them to communicate directly by reading and writing to shared memory.
 - Fast and uniform data sharing due to processors proximity to memory.
 - Global address leads to user friendly memory programming.
 - No scalability → more processors increases traffic on shared path.
 - Programmer responsibility for correct global memory access.
- Distributed Memory Model
 - Each processor has its own local memory, and communication between processors occurs via message-passing.
- Hybrid Model
 - Combines both shared and distributed memory models.



Parallel computing architectures: a taxonomy

The Flynn's taxonomy

- Based on
 - Multiple/Single Instruction (MI/SI)
 - Multiple/Single Data (MD/SD)
- Main architectural classes:
 - Single Instruction Single Data (SISD)
 - Single Instruction Multiple Data (SIMD)
 - Multiple Instruction Single Data (MISD)
 - Multiple Instruction Multiple Data (MIMD)
- Johnson's extension of MIMD
 - Based on:
 - Shared variables/Message Passing (SV/MP)
 - Global memory/distributed memory (GM/DM)



Architectures for modern HPC systems

- Multicore Processors:
 - Multiple cores on a single processor chip, allowing for simultaneous execution of multiple threads.
 - Ex: Intel's multi-core CPUs and AMD's Ryzen processors.
- Manycore Coprocessors
 - feature hundreds to thousands of cores optimized for data-parallel operations.
 - used for massively parallel workloads, such as matrix multiplications and deep learning model training.
 - Examples include Intel's Xeon Phi and NVIDIA's GPUs.
- Clusters and Supercomputers
 - Large-scale parallel systems made up of many interconnected nodes, each with its own memory and processors.
 - Ex: Cray supercomputers, Google's data centers.
- GPU Architectures
 - GPUs are specialized processors designed for parallelism, with thousands of cores to handle many computations simultaneously.
 - Ex: NVIDIA GPUs, used for tasks like matrix multiplication, deep learning, and graphics rendering.

Challenges in parallel programming

- **Concurrency Issues**

- Race Conditions: Occur when multiple threads/processes access shared data simultaneously without proper synchronization.
- Deadlock: A condition where processes are waiting for each other indefinitely, resulting in no progress.

- **Load Balancing:**

- Distributing work evenly across processors to avoid underutilization and idling cores.
- Granularity: The size of the tasks assigned to each processor can impact efficiency.

- **Scalability:**

- As the number of processors increases, the overhead of managing them (communication, synchronization) also increases.

- **Memory Management:**

- Ensuring that memory is used efficiently in parallel applications is critical to prevent bottlenecks.

Parallel programming tools

- OpenMP:
 - A set of compiler directives for shared memory parallelism in C, C++, and Fortran.
 - Easier to use for simple parallelism in multi-core systems.
 - Syntax and usage in C/C++.
- MPI:
 - A standard for message passing in distributed memory systems.
 - Allows programs to communicate by sending and receiving messages.
- CUDA:
 - Programming model for NVIDIA GPUs.
 - Uses GPU cores to perform large-scale parallel computations, ideal for data-intensive tasks like matrix multiplication.
- OpenCL:
 - A cross-platform framework for parallel programming on CPUs, GPUs, and other processors.
 - Supports heterogeneous computing systems.

3. Parallel computing models



Parallel Computing Models

- Parallel algorithms are designed with an assumption of an architecture of a parallel computer.
 - Different models of machines have been assumed for parallel algorithm design
- Requirements for parallel models
 - **Simplicity:**
 - A parallel model should allow easy analysis of various performance measures (speed, communication, memory utilization, etc.).
 - Results should be as hardware-independent as possible.
 - **Implementability:**
 - Parallel algorithms developed in a model should be easily realizable on a parallel machine.
 - Theoretical analysis should carry over and give meaningful performance estimates.
- A real satisfactory model does not exist!
 - We will study: Computation Graph; PRAM and Network models (Lecture 2)

DAG model

Directed Acyclic Graphs (DAGs) (or computation graphs) are often used to automatically parallelize numerical computations:

- nodes represent operations (single instructions or larger blocks)
 - In-degree is at most 2
 - Nodes without incoming edges correspond to input data
- edges represent dependencies (precedence constraints)
- branching instructions cannot be modelled
- completely hardware independent
- scheduling is not defined




DAG Model

- The DAG itself is not a complete algorithm A scheduling implements the algorithm on a parallel machine by assigning a time-step, t_v and a processor p_v to every node.
- A scheduling of a DAG $G = (V, E)$ on p processors is an assignment of pairs (t_v, p_v) to every internal node $v \in V$ subject to:
 - $p_v \in \{1, \dots, p\}; t_v \in \{1, \dots, T\}$
 - $t_u = t_v \rightarrow p_u \neq p_v$
 - $(u, v) \in E \rightarrow t_v \geq t_u + 1$

Where a non-internal node x (an input node) has $t_x = 0$.

DAG Example

- Example: product of 8 numbers



- The task of multiplying 8 numbers - using 4 processors in 3 units of time
 - Complete binary tree with n nodes is of height $\log_2(n)$
 - Multiplying n numbers can be done with $n/2$ processors in $\log_2(n)$ time units.
- Task assignment to processors \rightarrow scheduling:
 $SCH(p) := (p, t)$

Schedule	Explanation
$SCH(1) = (1,1)$	Processor 1 performs sub-task 1 at time 1
$SCH(2) = (2,1)$	Processor 2 performs sub-task 2 at time 1
$SCH(3) = (3,1)$	Processor 3 performs sub-task 3 at time 1
$SCH(4) = (4,1)$	Processor 4 performs sub-task 4 at time 1
$SCH(5) = (1,2)$	Processor 1 performs sub-task 5 at time 2
$SCH(6) = (2,2)$	Processor 2 performs sub-task 6 at time 2
$SCH(7) = (1,3)$	Processor 1 performs sub-task 7 at time 3

Schedule length = 3

PRAM Model

In Parallel Random-Access Machine (PRAM)

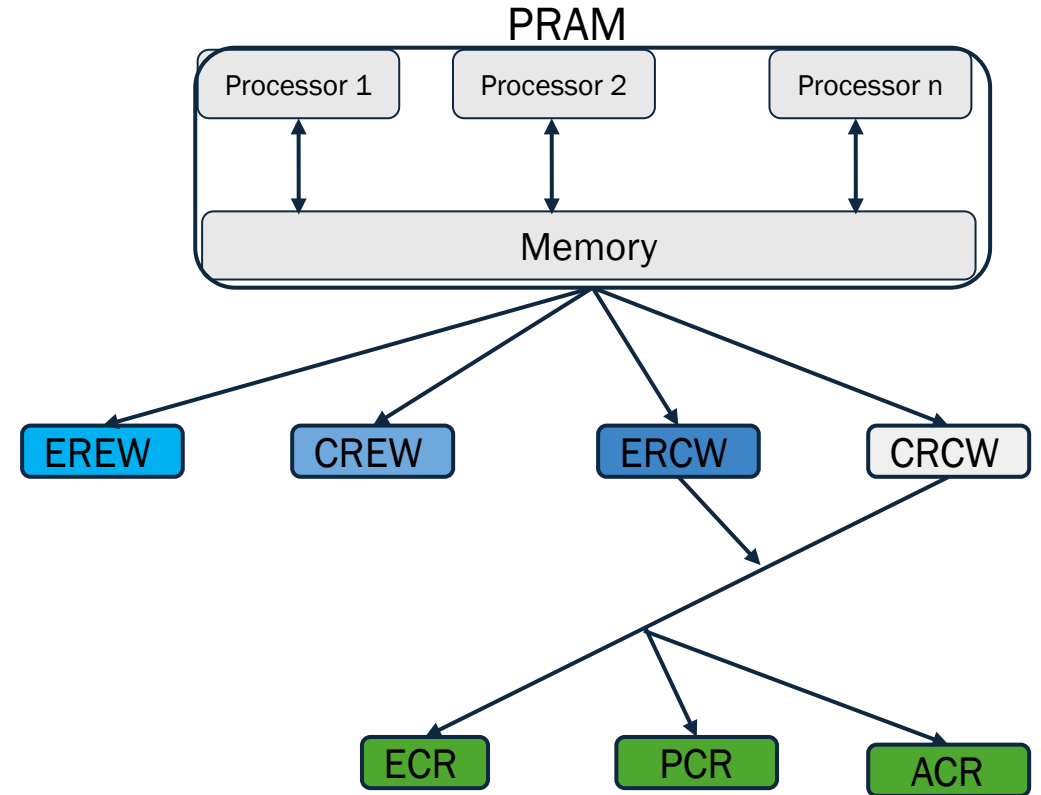
- also called **shared-memory model**
- all processors are connected in parallel to a global shared memory
 - all processors can read/write to memory
 - communications occur only via the shared memory
 - are synchronized via a common clock

PRAM model for SIMD architectures:

- Exclusive Read Exclusive Write (EREW)
- Concurrent Read Exclusive Write (CREW)
- Exclusive Read Concurrent Write (CRCW)
- Concurrent Read Concurrent Write (CRCW)

What happens if ≥ 2 processors attempt concurrent write?

- Conflict resolution, e.g.-
 - Equality/Common Conflict Resolution (ECR)
 - Priority Conflict Resolution (PCR)
 - Arbitrary/random Conflict Resolution (ACR)



PRAM – Simultaneous Read/Write Example

- Consider the following concurrent write operations on X by 3 processors
 $p_1(50 \rightarrow X), \quad p_2(60 \rightarrow X), \quad p_3(70 \rightarrow X)$

What is the outcome of the following (if p_1 has the highest priority)?

- Common CRCW or ERCW:
- Priority CRCW:
- Random CRCW:
- *Sum CRCW:

Parallel processing Effectiveness

Performance metrics

- Speed-up: $S_p(n) = \frac{T_1^*(n)}{T_p(n)}$
- Efficiency: $E_p(n) = \frac{T_1^*(n)}{pT_p(n)}$
- Redundancy: $R_p(n) = \frac{W_p(n)}{W_1(n)}$
- Utilization: $U_p(n) = \frac{W_p(n)}{pT(p)}$
- Communication overhead: (next lecture)

What is the significance of each of these metrics for a parallel program?

Performance metrics

- How are these parallel performance metrics related?
 - $1 \leq S_p(n) \leq p$
 - $U_p(n) = R_p(n)E_p(n)$
 - $E_p(n) = \frac{S_p(n)}{p}$
 - $\frac{1}{p} \leq E_p(n) \leq U_p(n) \leq 1$
 - $1 \leq R_p(n) \leq \frac{1}{E_p(n)} \leq p$

Amdahl's Law and Scalability

- **Amdahl's Law**

- Describes the theoretical maximum speedup of a program when only part of the program can be parallelized.

$$\text{Speedup} = \frac{1}{(1 - \alpha) + \frac{\alpha}{p}}$$

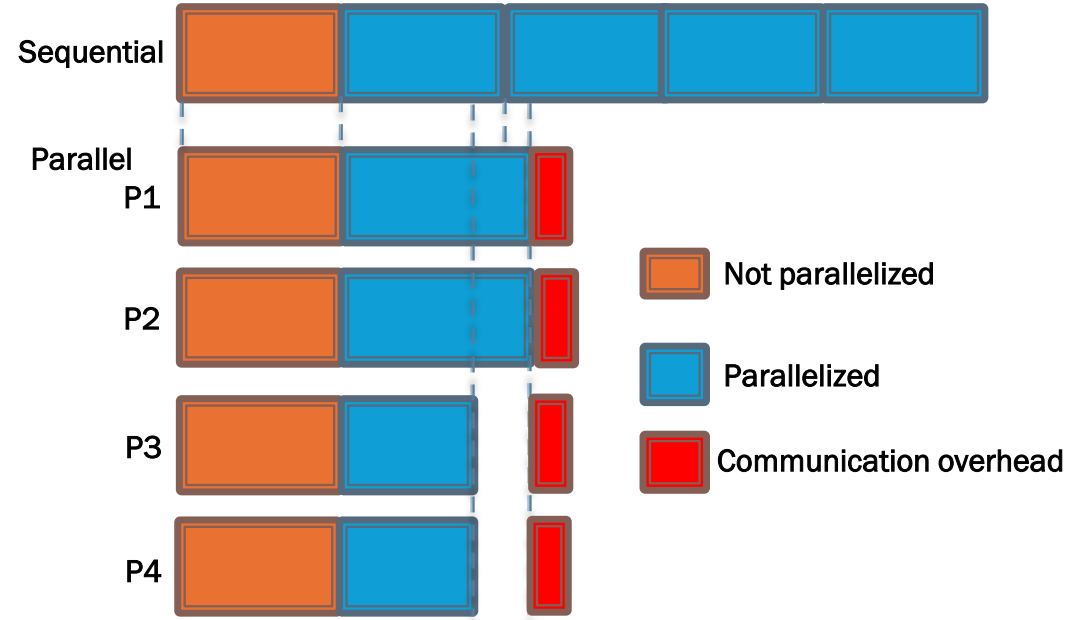

- α : proportion of the program that can be parallelized.
- p : number of processors.
- Even if you have many processors, the **sequential part of the program will limit the overall speedup**.

- **Scalability:**

- The ability of a parallel program to effectively utilize additional resources.
- Superlinear scalability can occur if communication overhead decreases with more processors.

Amdahl's Law – examples

- Execution of Program A using $p = 4$ processors
 - 80% can be parallelized \rightarrow 20% cannot be parallelized
- Parallel running time: $(1 - 0.8) + \frac{0.8}{4} = 0.4 \rightarrow 40\%$ of the sequential execution time.



Example: Sum of 16 numbers with 8 processors

- Example: Sum of 16 numbers with $p = 8$.

- Assuming unit-time additions and **ignoring all overheads:**


$$W_8(16) = \dots\dots\dots;$$

$$T_8(16) = \dots\dots$$

$$E_8(16) = \dots\dots\dots$$

$$S_8(16) = \dots\dots\dots$$

$$R_8(16) = \dots\dots\dots$$



Example: Sum of 16 numbers with 8 processors

- New assumption:
 - Vertically aligned operations are performed by the same processor
 - Each interprocessor transfer requires one unit of work (time)
- Performance:

$$W_8(16) = 22$$


$$T_8(16) = 7$$

$$E_8(16) = 15 / (8 \times 7) = 0.27 \Rightarrow 27\%$$

$$S_8(16) = 15 / 7 = 2.14$$

$$R_8(16) = 22 / 15 = 1.47$$

- Why the difference? Which solution is more practical?



Summary and Exercises

Summary

- Today, single-core performance is improving very slowly
 - To run programs significantly faster, programs must utilize multiple processing elements
 - you need to know how to design parallel algorithms and write parallel code
- Developing parallel algorithms and writing parallel programs can be challenging
 - Requires problem partitioning, communication, synchronization
 - Knowledge of machine characteristics is important
 - Performance may vary significantly from machine to machine
- Different models exist for studying parallel machine characteristics
 - Computation graph
 - Network models - Chain, ring, mesh, hypercubes, etc (next lecture)
 - PRAM - global shared memory model
 - unrealistic in practice due to lack of uniform memory access; global synchronization near impossible
 - good for understanding parallel techniques but not for algorithm development
- Measuring performance of parallel programs:
 - Speedup, efficiency, redundancy, utilization, communication cost (next lecture)

Exercises

- Exercise problems are available on Moodle.